

From Preferences to Values: Evaluating Latent User Understanding and Transfer in LLMs

Taehyung Noh
Hanyang University
Seoul, Republic of Korea
yestaehyung@hanyang.ac.kr

Beejin Son
Hanyang University
Seoul, Republic of Korea
beejinson@hanyang.ac.kr

Haein Yeo
Hanyang University
Seoul, Republic of Korea
haeinyeo@hanyang.ac.kr

Kyungsik Han*
Hanyang University
Seoul, Republic of Korea
kyungsikhan@hanyang.ac.kr

Abstract

Personalized LLM agents should adapt to users who often reveal preferences through behavioral cues rather than explicit statements. Because existing benchmarks mainly evaluate surface-level profile matching, we propose **PALETTE** 🎨, a theory-driven benchmark designed to evaluate value-based user understanding in LLMs. Grounded in Schwartz’s ten basic values, PALETTE contains 25,200 dialogues across 400 personas, pairing a historical dialogue that reveals a situational preference with a new context request that necessitates the same underlying value. We evaluate models on *Dialogue-Based Understanding* and *New-Context Application* to examine whether they can infer latent values from behavioral signals and generalize them across contexts. Our experiments with seven LLMs demonstrate that while models generate contextually appropriate dialog responses, they struggle to project an understanding of hidden values onto new contexts. PALETTE thus provides a diagnostic framework for verifying whether personalized agents can move beyond mere preference matching toward genuine, value-based user understanding. We make **PALETTE** publicly available at <https://taehyungnoh.com/cikm26-palette/>.

CCS Concepts

• **Human-centered computing** → **User models**; • **Computing methodologies** → **Natural language generation**.

Keywords

Personalization, User Understanding, User Preferences, User Values

ACM Reference Format:

Taehyung Noh, Haein Yeo, Beejin Son, and Kyungsik Han. 2026. From Preferences to Values: Evaluating Latent User Understanding and Transfer in LLMs. In *Proceedings of the 35th ACM International Conference on Information and Knowledge Management (CIKM '26)*, November 7–11, 2026, Rome, Italy. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3799682.3840043>

*Corresponding author



This work is licensed under a Creative Commons Attribution 4.0 International License. *CIKM '26, Rome, Italy*.

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2539-5/2026/11

<https://doi.org/10.1145/3799682.3840043>

1 Introduction

With the rapid advancement of Large Language Models (LLMs), personalized AI agents are transitioning from simple task-executors to long-term digital companions [9]. To provide seamless and proactive assistance, these agents should align with users over repeated interactions [17]. However, in real-world, rather than providing explicit, structured profiles of their desires, users tend to reveal their tendencies through subtle behavioral cues, such as correcting an agent’s output or interrupting long model generations to demand immediate conclusions [8, 21]. Consequently, an intelligent personalized agent should be able to read between the lines, inferring the latent driving forces behind a user’s explicit actions.

However, existing personalization benchmarks appear to focus on surface-level preference following or explicit profile matching [18]. These frameworks evaluate whether an LLM can replicate observed patterns within the same or similar contexts [6]. However, a fundamental limitation remains overlooked: surface preferences are highly situational, whereas the human decision-making process is governed by more profound, stable drivers known as values. According to social psychology, values are trans-situational goals that serve as guiding principles in people’s lives, and the same surface preference can stem from entirely different latent values [1]. For instance, a user’s preference for a “short, concise email summary” might reflect Achievement (prioritizing efficient task completion), Self-Direction (retaining autonomy over which details to inspect), or Security (reducing uncertainty and avoiding errors). If an agent merely memorizes the surface preference without grasping the underlying value, it will inevitably fail to adapt when the user moves to a new context, where the exact wording cues of the preference disappear but the core value remains identical.

To bridge this gap, we introduce **PALETTE** 🎨, a novel evaluation benchmark designed to examine LLMs’ capacity for value-based user understanding and cross-context transfer. Grounded in Schwartz’s Theory of Basic Human Values [16], PALETTE shifts the evaluation paradigm from superficial behavior replication to latent motive inference. We formalize this via two distinct, sequential tasks: (1) *Dialogue-Based Understanding*, which tests whether an LLM can accurately infer a user’s latent value solely from behavioral and conversational cues, and (2) *New-Context Application*, which evaluates if the model can successfully project that inferred value into a completely unseen, wordingly decoupled domain to satisfy the user’s implicit expectations.




Figure 1: Overview of the PALETTE evaluation framework. Models infer a user’s hidden value and situational preference from implicit behavioral cues during an ongoing dialogue (Task 1), and subsequently apply this inferred understanding to generate appropriate responses in a topically distinct new context (Task 2).

To ensure the rigorousness of PALETTE, we constructed a dataset comprising 25,200 dialogues derived from 400 distinct, multi-value user personas. A key technical contribution of our data pipeline lies in the explicit decoupling and multi-faceted abstraction of user contexts—*latent values*, *situational preferences*, and *scenario environments*. This structural separation prevents LLMs from exploiting superficial text-matching shortcuts, forcing them to rely on genuine abstract reasoning and cross-domain value transfer. Furthermore, we decouple evaluation metrics into *inference accuracy* and *response appropriateness* to dissect whether a model truly comprehends the user’s underlying motivations or simply generates plausible-sounding dialogue by chance.

Through extensive benchmarking of seven LLMs, we identify a dichotomy in current personalization architectures. While modern LLMs are remarkably adept at maintaining conversational fluency and surface appropriateness, they struggle with true latent inference and cross-context value transfer. In particular, response appropriateness drops by up to 25.1 % points when models are transitioned into decoupled domains, exposing a critical vulnerability in current agent alignment paradigms.

In summary, our main contributions are as follows:

- We redefine the personalization problem by distinguishing surface-level situational preferences from trans-situational latent values, introducing the concept of *value transfer* to the LLM evaluation literature.
- We develop a rigorously decoupled dataset of 25,200 dialogues grounded in social science theory to benchmark value-driven user understanding without shortcut exploitation.
- We conduct a comprehensive audit of multiple LLMs, revealing a critical gap between conversational appropriateness and latent value comprehension, offering actionable directions for future personalized agent architectures.

2 Background

Recent studies have evaluated whether LLMs can infer user preferences from dialogue contexts, user histories, or preference feedback and reflect them in later responses. For example, PrefEval [22] evaluates preference-following capabilities within long multi-session dialogue contexts, while RealPref [5] studies diverse preference expressions in long-term interactions and generalization to unseen

scenarios. ALOE [21] evaluates preference-aligned response generation through persona-conditioned multi-turn interactions, and PersonalLLM [25] investigates personalization under user-specific preference distributions. CUPID [8] focuses on context-dependent preferences revealed through previous multi-turn feedback and separately evaluates preference inference and response generation. These studies show that inferring implicit preferences and applying them to subsequent or unseen requests are central challenges in LLM personalization. However, existing frameworks typically treat preferences as surface-level behavioral markers; they neither formalize the latent value underlying a preference as an explicit gold label, nor do they jointly control behaviorally revealed preferences, latent values, and new-context requests across divergent topics and surface expressions. Previous evaluations largely test whether models can recover a user’s revealed preference from observed choices. PALETTE shifts this paradigm by evaluating whether models can identify the core latent value expressed by these choices and leverage that value to respond appropriately to a new request.

3 PALETTE

3.1 Task Formulation

Each PALETTE scenario is formulated as $\langle p, v^*, \pi^*, s, h \rangle$: user persona p , latent value v^* , situational preference π^* , original setting s , and new context h . During evaluation, models observe task contexts but not the ground-truth v^* or π^* . We evaluate these scenarios via two sequential tasks.

In *Dialogue-Based Understanding*, the model interacts with a user agent in s , whose preferences are conveyed implicitly through behavioral cues. At each turn, the model generates a user-visible response and a hidden user-understanding prediction (inferring π^* and v^*). Separating v^* and π^* in this hidden prediction ensures we evaluate whether models infer the underlying value motivating a preference, rather than just recognizing the preference itself.

In *New-Context Application*, the model receives the s dialogue history and the new context h to generate a single-turn response and an updated inference. Because s and h are intentionally designed to be lexically and topically distinct, this task strictly assesses whether the model can adapt the inferred value to new situations without relying on surface matching.

Table 1: PALETTE dataset summary.

Item	Count	Description
<i>Controlled scenarios</i>		
Personas	400	User personas.
Values per persona	3	Assigned values.
Scenarios	1,200	Persona–value pairs.
<i>Generated dialogue records</i>		
Target models	7	Evaluated LLMs.
Dialogue records per model	3,600	Per-model records.
Total dialogue records	25,200	Full released records.

Table 2: Situation overlap between original and new contexts.

Measure	Paired	Random ref.	Diff.
Word overlap	2.4%	2.4%	−0.1
Embedding similarity	0.25	0.16	+0.09

3.2 Scenario Synthesis

We synthesize 25,200 PALETTE dialogues using 400 diverse personas from Nemotron-Personas-USA [12]. To decouple target values from demographic stereotypes, we independently assign each persona three distinct, balanced Schwartz values. For each assignment, we generate a situational preference π^* , alongside an original setting s and a new context h . Crucially, s and h are generated with distinct topics and language to ensure applying the preference in h requires value abstraction rather than surface-level text reuse.

To enable strict evaluation, we annotate scenarios with confounder values—plausible alternative values that could similarly motivate π^* . Finally, to prevent information leakage, we manually reviewed all candidates to check if the observable text directly revealed the ground-truth v^* or π^* . Based on this criterion, we dropped and regenerated 127 leaky scenarios to finalize the dataset.

3.3 Judge-based Scoring

Following previous work [8, 21, 22], PALETTE decouples inference and generation errors by reporting two metrics: understanding accuracy (UA) and response appropriateness (RA). For UA, an LLM-based judge evaluates the model’s hidden prediction against the ground-truth metadata. UA is marked correct only when the inferred value exactly matches the Schwartz value v^* to the exclusion of neighboring values, while the inferred preference semantically aligns with π^* . For RA, the judge assesses the user-facing response against v^* , π^* , a checklist, and preference keypoints, independent of the model’s inference. A response is deemed appropriate if it satisfies $\geq 50\%$ of both checklist items and keypoints. In *New-Context Application*, RA additionally penalizes application errors—defined as the improper replication of context-specific patterns from s instead of adapting the preference to h .

3.4 Dataset Summary

Table 1 summarizes the PALETTE, generated from 400 personas configured with three balanced Schwartz values each. Each scenario comprises v^* , π^* , s , h , a confounder value, and scoring criteria.

Table 3: Human validation results for scenario quality and value recoverability on 200 sampled scenarios.

Task	Metric	Result
Value identification (Ground truth hidden)	Mean individual accuracy	71.5%
	Majority-vote accuracy	73.0%
	Fleiss’ κ	0.796
Validity rating (Ground truth shown, 7-point Likert scale)	Value validity: mean (% of ratings ≥ 6)	6.43 / 94%
	Pref. naturalness: mean (% of ratings ≥ 6)	6.39 / 92%
	New-context validity: mean (% of ratings ≥ 6)	6.04 / 79%
	Pairwise agreement within one point	95.6%

Table 4: Judge validation against human annotations. Human labels are aggregated by majority vote.

Scoring decision	n	Human κ	Judge Acc.	Judge κ
UA: preference correctness	200	0.642	82.9%	0.465
UA: joint correctness	200	0.853	89.5%	0.790
RA: appropriateness	200	0.695	71.0%	0.420

4 Dataset Validation

4.1 Situation Overlap Analysis

In *New-Context Application*, h is specifically designed to diverge topically and lexically from s . This separation ensures that the task evaluates true preference transfer rather than surface-level pattern matching. To verify this design, we compute stop-word-filtered word overlap and embedding cosine similarity using `all-mpnet-base-v2` [15] between s and h . As shown in Table 2, the comparison against random pairs confirms our intended design: the scenarios successfully avoid surface-level continuity while remaining robustly anchored by the same underlying value-based preference.

4.2 Scenario Validation

We validate whether the synthesized scenarios reflect PALETTE’s intended value-theoretic structure. Three human annotators evaluated 200 sampled scenarios (20 per Schwartz value) across two tasks. First, in a value-identification task with the ground-truth hidden, annotators chose the most appropriate value among the ten Schwartz values based on the persona description, s , and behavioral evidence. Second, with the ground-truth revealed, annotators rated three dimensions on a 7-point Likert scale: (i) value validity (given the persona and behavior), (ii) preference naturalness, and (iii) new-context validity (whether h requires the same value-grounded criterion despite differing from s in surface form).

Table 3 reports the results. Annotators recovered the intended value with 71.5% mean individual accuracy and 73.0% majority-vote accuracy, with a Fleiss’s κ of 0.796, indicating inter-annotator agreement and proving that the context provides highly recoverable evidence for the latent value. Furthermore, rating results show the scenarios are highly well-formed: value validity (mean 6.43, $94\% \geq 6$) and preference naturalness (mean 6.39, $92\% \geq 6$) received strong positive ratings. New-context validity scored slightly lower (mean 6.04, $79\% \geq 6$), reflecting the difficulty of decoupling h from s while preserving the preference criterion.

Table 5: PALETTE main results (mean $\pm$ SD) for UA and RA. RA-UA measures the inference-response discrepancy within each task. Δ RA (New-context RA – Dialogue RA) measures the retention of response appropriateness across contexts. **Blue bold: best performance or transfer stability; **red bold**: largest discrepancy or degradation.**

Model	Dialogue-Based Understanding			New-Context Application			Transfer
	UA $\uparrow$	RA $\uparrow$	RA-UA	UA $\uparrow$	RA $\uparrow$	RA-UA	Δ RA $\uparrow$
Llama-3.1-8B [4]	32.7 $\pm$ 2.1	51.4 $\pm$ 2.7	+18.7	32.9 $\pm$ 2.3	26.3 $\pm$ 1.9	-6.6	-25.1
Gpt-oss-20B [14]	41.6 $\pm$ 2.3	69.9$\pm$2.7	+28.3	52.6 $\pm$ 2.2	53.8 $\pm$ 2.4	+1.2	-16.1
Gemma-3-27B [19]	48.6 $\pm$ 2.1	66.1 $\pm$ 2.3	+17.5	54.4 $\pm$ 2.9	54.4 $\pm$ 2.3	0.0	-11.7
GLM-4-32B [3]	30.5 $\pm$ 1.6	42.4 $\pm$ 2.9	+11.9	37.0 $\pm$ 2.9	26.2 $\pm$ 2.0	-10.8	-16.2
Llama-3.3-70B [11]	28.5 $\pm$ 3.9	45.0 $\pm$ 4.0	+16.5	33.4 $\pm$ 3.3	26.5 $\pm$ 1.2	-6.9	-18.5
DeepSeek-V3.1 [2]	48.9$\pm$2.8	62.1 $\pm$ 2.6	+13.2	62.1$\pm$1.8	59.2$\pm$2.2	-2.9	-2.9
Gpt-4o-mini [13]	34.7 $\pm$ 0.5	46.7 $\pm$ 0.2	+12.0	39.2 $\pm$ 0.6	32.7 $\pm$ 1.0	-6.5	-14.0

4.3 Judge Validation

Following previous work [7, 10, 24], we validate PALETTE’s LLM judge against a human-majority baseline. Three annotators evaluated a held-out set of model outputs using the same scoring rubric as the LLM judge. We report three scoring targets: UA preference correctness, UA joint correctness, and RA appropriateness. UA preference correctness evaluates whether the inferred situational preference matches the target preference. UA joint correctness follows the PALETTE definition, marked as correct only when both the inferred value and the inferred preference are correct. RA appropriateness evaluates whether the user-facing response satisfies the target preference based on checklist and keypoint coverage.

Table 4 reports the results. The judge achieved the strongest agreement on UA joint correctness, with 89.5% accuracy and a Cohen’s κ of 0.790, closely trailing the high baseline human–human agreement ($\kappa = 0.853$). UA preference correctness also showed high judge accuracy (82.9%), although its chance-corrected agreement was more moderate ($\kappa = 0.465$). RA appropriateness was comparatively more difficult, yielding an accuracy of 71.0% and a Cohen’s κ of 0.420, reflecting the open-ended nature of response-level evaluation. Overall, these results justify the LLM judge for aggregate comparisons, while individual response-level judgments warrant caution.

5 Experiments

5.1 Experimental Setup

We evaluate seven LLMs: Llama-3.1-8B, Gpt-oss-20B, Gemma-3-27B, GLM-4-32B, DeepSeek-V3.1, Llama-3.3-70B, and Gpt-4o-mini. Following prior personalization benchmarks [23], we use Qwen-3-235B [20] as a shared user ($T=0.3$) and judge ($T=0.0$) agent. We report UA and RA (§3.3) on both tasks. Each scenario is run three times per model to capture dialogue variation, yielding 3,600 dialogues per model, for which we report mean and standard deviation.

5.2 Latent Value Inference and New-Context Application

Table 5 reports UA and RA across both tasks. In the *Dialogue-Based Understanding* setting, RA consistently exceeds UA, reaching 70% for RA compared to 49% for UA. This gap suggests that models can generate superficially appropriate dialogue responses without accurately inferring the user’s latent values and preferences.

In *New-Context Application*, models show a clear performance split in their ability to apply inferred values onto new contexts. DeepSeek-V3.1, Gemma-3-27B, and Gpt-oss-20B achieve scores ranging from 53.8 to 59.2% in RA, whereas Llama-3.1-8B, GLM-4-32B, and Llama-3.3-70B remain tightly clustered in a lower tier between 26.2 and 26.5%. GPT-4o-mini occupies a moderate middle ground at 32.7%. Notably, Gpt-oss-20B and Gemma-3-27B substantially outperform the larger Llama-3.3-70B in New-Context RA, suggesting that PALETTE captures differences in value-level user understanding that are not explained by model scale. To verify that *New-Context Application* requires information transferred from the original dialogue rather than relying on generic response generation, we add a no-history control that removes the dialogue history D_s and presents only the new request h . Incorporating the full context environment improves New-Context RA by 17.9 percentage points over the no-history baseline, confirming that historical dialogue contributes beyond answering the new request alone.

6 Conclusion

We introduced PALETTE, a theory-driven benchmark that evaluates whether LLMs infer latent values from behaviorally expressed preferences and transfer them to new contexts, decoupling understanding accuracy (UA) from response appropriateness (RA). Across seven LLMs, models generate superficially appropriate responses without accurately inferring hidden values, and their ability to transfer this understanding diverges substantially. A limitation is the reliance on an LLM-based judge for open-ended responses; our validation supports aggregate comparisons, but lower RA agreement means RA should be interpreted only at the model-comparison level. Overall, PALETTE surfaces alignment failures that preference satisfaction alone overlooks, offering a foundation for future value-based user-modeling evaluation. In future work, we plan to expand the benchmark with more diverse personas and scenarios and to investigate specialized tuning methods for effective value transfer.

Acknowledgments

This work was supported by IITP (No. RS-2020-II201373; No. RS-2025-25422680) and KIAT (No. RS-2025-02633871) grants funded by the Korea government.

7 GenAI Usage Disclosure

Generative AI models were used to generate the synthetic dialogue records and to run the shared-user and judge-agent components described in Sections 3 and 5.1. The authors also used generative AI tools for grammatical review and minor wording refinements of the manuscript. All research design, scenario criteria, annotation decisions, analysis, interpretation, and final manuscript decisions were made by the authors.

References

- [1] Anat Bardi and Shalom H Schwartz. 2003. Values and behavior: Strength and structure of relations. *Personality and social psychology bulletin* 29, 10 (2003), 1207–1220.
- [2] DeepSeek-AI. 2024. DeepSeek-V3 Technical Report. arXiv:2412.19437 [cs.CL] <https://arxiv.org/abs/2412.19437>
- [3] Team Glm, Aohan Zeng, Bin Xu, Bowen Wang, Chenhui Zhang, Da Yin, Dan Zhang, Diego Rojas, Guanyu Feng, Hanlin Zhao, et al. 2024. Chatglm: A family of large language models from glm-130b to glm-4 all tools. *arXiv preprint arXiv:2406.12793* (2024).
- [4] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, et al. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783* (2024).
- [5] Qianyun Guo, Yibo Li, Yue Liu, and Bryan Hooi. 2026. Towards Realistic Personalization: Evaluating Long-Horizon Preference Following in Personalized User-LLM Interactions. *arXiv preprint arXiv:2603.04191* (2026).
- [6] Bowen Jiang, Zhuoqun Hao, Young-Min Cho, Bryan Li, Yuan Yuan, Sihao Chen, Lyle Ungar, Camillo J Taylor, and Dan Roth. 2025. Know me, respond to me: Benchmarking llms for dynamic user profiling and personalized responses at scale. *arXiv preprint arXiv:2504.14225* (2025).
- [7] Seungone Kim, Jay Shin, Joel Jang, Shayne Longpre, Hwaran Lee, Sangdoo Yun, Ryan Shin, Sungdong Kim, James Thorne, Minjoon Seo, et al. 2024. Prometheus: Inducing fine-grained evaluation capability in language models. In *International Conference on Learning Representations*, Vol. 2024. 29927–29962.
- [8] Tae Soo Kim, Yoonjoo Lee, Yoonah Park, Jiho Kim, Young-Ho Kim, and Juho Kim. 2025. Cupid: Evaluating personalized and contextualized alignment of llms from interactions. *arXiv preprint arXiv:2508.01674* (2025).
- [9] Hao Li, Chenghao Yang, An Zhang, Yang Deng, Xiang Wang, and Tat-Seng Chua. 2025. Hello again! llm-powered personalized agent for long-term dialogue. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. 5259–5276.
- [10] Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. 2023. G-eval: NLG evaluation using gpt-4 with better human alignment. In *Proceedings of the 2023 conference on empirical methods in natural language processing*. 2511–2522.
- [11] Meta. 2024. Llama 3.3 70B Instruct. <https://huggingface.co/meta-llama/Llama-3.3-70B-Instruct>
- [12] NVIDIA. 2025. Nemotron-Personas-USA. (2025). <https://huggingface.co/datasets/nvidia/Nemotron-Personas-USA>
- [13] OpenAI. 2024. GPT-4o System Card. arXiv:2410.21276 [cs.CL]
- [14] OpenAI. 2025. gpt-oss-120b & gpt-oss-20b Model Card. arXiv:2508.10925 [cs.CL] <https://arxiv.org/abs/2508.10925>
- [15] Nils Reimers and Iryna Gurevych. 2019. Sentence-bert: Sentence embeddings using siamese bert-networks. In *Proceedings of the 2019 conference on empirical methods in natural language processing and the 9th international joint conference on natural language processing (EMNLP-IJCNLP)*. 3982–3992.
- [16] Shalom H Schwartz. 1992. Universals in the content and structure of values: Theoretical advances and empirical tests in 20 countries. In *Advances in experimental social psychology*. Vol. 25. Elsevier, 1–65.
- [17] Zhen Tan, Jun Yan, I-Hung Hsu, Rujun Han, Zifeng Wang, Long Le, Yiwen Song, Yanfei Chen, Hamid Palangi, George Lee, et al. 2025. In prospect and retrospect: Reflective memory management for long-term personalized dialogue agents. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 8416–8439.
- [18] Meiling Tao, Chenghao Zhu, Dongyi Ding, Tiannan Wang, Yuchen Eleanor Jiang, and Wangchunshu Zhou. 2025. Personafedback: A large-scale human-annotated benchmark for personalization. *arXiv preprint arXiv:2506.12915* (2025).
- [19] Gemma Team. 2025. Gemma 3. (2025). <https://goo.gle/Gemma3Report>
- [20] Qwen Team. 2025. Qwen3 Technical Report. arXiv:2505.09388 [cs.CL] <https://arxiv.org/abs/2505.09388>
- [21] Shujin Wu, Yi R Fung, Cheng Qian, Jeonghwan Kim, Dilek Hakkani-Tur, and Heng Ji. 2025. Aligning llms with individual preferences via interaction. In *Proceedings of the 31st International Conference on Computational Linguistics*. 7648–7662.
- [22] Siyan Zhao, Mingyi Hong, Yang Liu, Devamanyu Hazarika, and Kaixiang Lin. 2025. DO LLMS RECOGNIZE YOUR PREFERENCES? EVALUATING PERSONALIZED PREFERENCE FOLLOWING IN LLMS. In *13th International Conference on Learning Representations, ICLR 2025*. International Conference on Learning Representations, ICLR, 72531–72574.
- [23] Zheng Zhao, Clara Vania, Subhradeep Kayal, Naila Khan, Shay B Cohen, and Emine Yilmaz. 2025. Personalens: A benchmark for personalization evaluation in conversational ai assistants. In *Findings of the Association for Computational Linguistics: ACL 2025*. 18023–18055.
- [24] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. 2023. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in neural information processing systems* 36 (2023), 46595–46623.
- [25] Thomas Zollo, Andrew Siah, Naimeng Ye, Li Li, and Hongseok Namkoong. 2025. Personalllm: Tailoring llms to individual preferences. In *International Conference on Learning Representations*, Vol. 2025. 66949–66971.